

Intro to Econometrics: Recitation 8

Gustavo Pereira

November 12, 2019

Roadmap for today

- Chapter 6 quick review:
 - ▶ Asymptotic order notation
 - ▶ The basic toolkit: CMT, delta method and limit theorems
- BH Chapter 7:
 - ▶ Asymptotic properties in the LP/LR model

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

- ▶ How do we make this rigorous?

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

► How do we make this rigorous?

- “ $|a_n - b_n|$ gets large” is not a good idea; e.g. $a_n = 5n$, $b_n = n$

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

► How do we make this rigorous?

- “ $|a_n - b_n|$ gets large” is not a good idea; e.g. $a_n = 5n$, $b_n = n$
- Let $r_n = \frac{a_n}{b_n}$:

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

- ▶ How do we make this rigorous?
- “ $|a_n - b_n|$ gets large” is not a good idea; e.g. $a_n = 5n$, $b_n = n$
- Let $r_n = \frac{a_n}{b_n}$:
 - ▶ $a_n \in O(b_n) \iff |r_n|$ is bounded

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

- ▶ How do we make this rigorous?
- “ $|a_n - b_n|$ gets large” is not a good idea; e.g. $a_n = 5n$, $b_n = n$
- Let $r_n = \frac{a_n}{b_n}$:
 - ▶ $a_n \in O(b_n) \iff |r_n|$ is bounded
 - ▶ $a_n \in o(b_n) \iff |r_n| \rightarrow 0$

Asymptotic order notation

Deterministic sequences

- Consider the claim:

“The sequence $a_n = n^2$ diverges to infinity an order of magnitude faster than $b_n = n$ ”

- ▶ How do we make this rigorous?
- “ $|a_n - b_n|$ gets large” is not a good idea; e.g. $a_n = 5n$, $b_n = n$
- Let $r_n = \frac{a_n}{b_n}$:
 - ▶ $a_n \in O(b_n) \iff |r_n|$ is bounded
 - ▶ $a_n \in o(b_n) \iff |r_n| \rightarrow 0$
- Note: we’re assuming for simplicity that $b_n > 0$

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:

❶ $n \in o(n^2)$

❷ $\log(n) \in O(n)$

❸ $\sin(n) \in O(1)$

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:
 - 1 $n \in o(n^2)$
 - 2 $\log(n) \in O(n)$
 - 3 $\sin(n) \in O(1)$
- Also, for arbitrary (positive) x_n and y_n ,

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:

- ❶ $n \in o(n^2)$

- ❷ $\log(n) \in O(n)$

- ❸ $\sin(n) \in O(1)$

- Also, for arbitrary (positive) x_n and y_n ,

- ❶ $x_n \in O(x_n)$ but the same doesn't hold for $o(x_n)$

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:
 - ❶ $n \in o(n^2)$
 - ❷ $\log(n) \in O(n)$
 - ❸ $\sin(n) \in O(1)$
- Also, for arbitrary (positive) x_n and y_n ,
 - ❶ $x_n \in O(x_n)$ but the same doesn't hold for $o(x_n)$
 - ❷ If $x_n \in O(An^2 + Bn + C)$, then $x_n \in O(n^2)$; moreover, $x_n \in o(n^3)$
 - ★ Similar fact for higher degree polynomials goes through

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:
 - 1 $n \in o(n^2)$
 - 2 $\log(n) \in O(n)$
 - 3 $\sin(n) \in O(1)$
- Also, for arbitrary (positive) x_n and y_n ,
 - 1 $x_n \in O(x_n)$ but the same doesn't hold for $o(x_n)$
 - 2 If $x_n \in O(An^2 + Bn + C)$, then $x_n \in O(n^2)$; moreover, $x_n \in o(n^3)$
 - ★ Similar fact for higher degree polynomials goes through
 - 3 $x_n \in O(y_n)$ iff $x_n/y_n \in o(1)$; same for O

Asymptotic order notation

Deterministic sequences: examples

- Convince yourself (or even better, give a proof) that:
 - 1 $n \in o(n^2)$
 - 2 $\log(n) \in O(n)$
 - 3 $\sin(n) \in O(1)$
- Also, for arbitrary (positive) x_n and y_n ,
 - 1 $x_n \in O(x_n)$ but the same doesn't hold for $o(x_n)$
 - 2 If $x_n \in O(An^2 + Bn + C)$, then $x_n \in O(n^2)$; moreover, $x_n \in o(n^3)$
 - ★ Similar fact for higher degree polynomials goes through
 - 3 $x_n \in O(y_n)$ iff $x_n/y_n \in o(1)$; same for O
 - 4 If $x_n \rightarrow x$, then $x_n \in O(1)$

Asymptotic order notation

Deterministic sequences: results

Lemma

For any positive sequences x_n , y_n , z_n and w_n ,

- ❶ If $x_n \in o(z_n)$ and $y_n \in o(w_n)$, then

$$x_n y_n \in o(z_n w_n)$$

$$x_n + y_n \in o(z_n + w_n)$$

and the same holds for O

- ❷ If $x_n \in o(z_n)$ and $y_n \in O(w_n)$, then

$$x_n y_n \in o(z_n w_n)$$

$$x_n + y_n \in O(z_n + w_n)$$

Asymptotic order notation

Deterministic sequences: results

- Proof:

- ▶ First part of (1) follows from limit arithmetic
- ▶ First part of (2) is a consequence of the fact that $a_n \rightarrow 0$ and b_n bounded implies $a_n b_n \rightarrow 0$
- ▶ Second part of (1) and (2): use

$$\frac{x_n + y_n}{z_n + w_n} = \frac{x_n}{z_n} \frac{z_n}{z_n + w_n} + \frac{y_n}{w_n} \frac{w_n}{z_n + w_n}$$

the ratio $z_n/(z_n + w_n) \in [0, 1]$ is bounded

Asymptotic order notation

Sequences of random variables

- Analogs: let X_n be a sequence of random variables and a_n a positive deterministic sequence
 - ▶ $X_n \in O_p(a_n)$ iff $\frac{X_n}{a_n}$ is *stochastically bounded*
 - ▶ $X_n \in o_p(a_n)$ iff $\frac{X_n}{a_n} \xrightarrow{p} 0$

Asymptotic order notation

Sequences of random variables

- Analogs: let X_n be a sequence of random variables and a_n a positive deterministic sequence
 - ▶ $X_n \in O_p(a_n)$ iff $\frac{X_n}{a_n}$ is *stochastically bounded*
 - ▶ $X_n \in o_p(a_n)$ iff $\frac{X_n}{a_n} \xrightarrow{p} 0$
- Definition of stochastically bounded.

For every $\delta > 0$ there exists M_δ such that

$$\Pr\{|X_n| \leq M_\delta\} > 1 - \delta$$

Asymptotic order notation

Sequences of random variables

- Analogs: let X_n be a sequence of random variables and a_n a positive deterministic sequence

- ▶ $X_n \in O_p(a_n)$ iff $\frac{X_n}{a_n}$ is *stochastically bounded*

- ▶ $X_n \in o_p(a_n)$ iff $\frac{X_n}{a_n} \xrightarrow{p} 0$

- Definition of stochastically bounded.

For every $\delta > 0$ there exists M_δ such that

$$\Pr\{|X_n| \leq M_\delta\} > 1 - \delta$$

- Result: lemma goes through with o and O replaced with o_p and O_p

Asymptotic order notation

Sequences of random variables: an useful result

A sequence (X_n) that converges in distribution is necessarily $O_p(1)$

Asymptotic order notation

Sequences of random variables: an useful result

A sequence (X_n) that converges in distribution is necessarily $O_p(1)$

- Proof. Let Y be st $X_n \xrightarrow{d} Y$. Denote F_Y and F_n the CDF of Y and X_n , respectively. Fix $\delta > 0$.

- ▶ Let M_δ be such that

$$F_Y(M_\delta) - F_Y(-M_\delta) = \Pr\{Y \in (-M_\delta, M_\delta]\} > 1 - \frac{\delta}{2}$$

- ▶ In addition, we can choose M_δ st F_Y continuous at M_δ and $-M_\delta$ (why?)

Asymptotic order notation

Sequences of random variables: an useful result

A sequence (X_n) that converges in distribution is necessarily $O_p(1)$

- Proof. Let Y be st $X_n \xrightarrow{d} Y$. Denote F_Y and F_n the CDF of Y and X_n , respectively. Fix $\delta > 0$.

► Let M_δ be such that

$$F_Y(M_\delta) - F_Y(-M_\delta) = \Pr\{Y \in (-M_\delta, M_\delta]\} > 1 - \frac{\delta}{2}$$

- In addition, we can choose M_δ st F_Y continuous at M_δ and $-M_\delta$ (why?)
- Convergence in distribution implies $F_n(M_\delta) \rightarrow F_Y(M_\delta)$ and $F_n(-M_\delta) \rightarrow F_Y(-M_\delta)$

Asymptotic order notation

Sequences of random variables: an useful result

A sequence (X_n) that converges in distribution is necessarily $O_p(1)$

- Proof. Let Y be st $X_n \xrightarrow{d} Y$. Denote F_Y and F_n the CDF of Y and X_n , respectively. Fix $\delta > 0$.

► Let M_δ be such that

$$F_Y(M_\delta) - F_Y(-M_\delta) = \Pr\{Y \in (-M_\delta, M_\delta]\} > 1 - \frac{\delta}{2}$$

- In addition, we can choose M_δ st F_Y continuous at M_δ and $-M_\delta$ (why?)
- Convergence in distribution implies $F_n(M_\delta) \rightarrow F_Y(M_\delta)$ and $F_n(-M_\delta) \rightarrow F_Y(-M_\delta)$
- Hence for n_0 sufficiently large, $F_n(M_\delta) > F_Y(M_\delta) - \delta/4$ and $F_n(-M_\delta) < F_Y(-M_\delta) + \delta/4$ for all $n \geq n_0$

Asymptotic order notation

Sequences of random variables: an useful result

A sequence (X_n) that converges in distribution is necessarily $O_p(1)$

- Proof. Let Y be st $X_n \xrightarrow{d} Y$. Denote F_Y and F_n the CDF of Y and X_n , respectively. Fix $\delta > 0$.

► Let M_δ be such that

$$F_Y(M_\delta) - F_Y(-M_\delta) = \Pr\{Y \in (-M_\delta, M_\delta]\} > 1 - \frac{\delta}{2}$$

► In addition, we can choose M_δ st F_Y continuous at M_δ and $-M_\delta$ (why?)

- Convergence in distribution implies $F_n(M_\delta) \rightarrow F_Y(M_\delta)$ and $F_n(-M_\delta) \rightarrow F_Y(-M_\delta)$
- Hence for n_0 sufficiently large, $F_n(M_\delta) > F_Y(M_\delta) - \delta/4$ and $F_n(-M_\delta) < F_Y(-M_\delta) + \delta/4$ for all $n \geq n_0$
- That implies

$$\Pr\{|X_n| \leq M_\delta\} \geq F_n(M_\delta) - F_n(-M_\delta) > 1 - \delta$$

The asymptotics toolkit

- You need to know the following by heart:

- ① General results

- ★ Continuous mapping theorem
 - ★ Slutsky's theorem
 - ★ Delta method

- ② Results about averages

- ★ Law of large numbers
 - ★ Central Limit Theorem

Continuous mapping theorem

- Let $F : U \rightarrow V$ have a set of discontinuities U_0
- If $X_n \xrightarrow{d} X$ or $X_n \xrightarrow{p} X$, then $F(X_n) \rightarrow F(X)$ as long as

$$\Pr\{X \in U_0\} = 0$$

Slutsky's theorem

- If $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{p} c$ where c is non-random, then

- ❶ $X_n + Y_n \xrightarrow{d} X + c$

- ❷ $X_n Y_n \xrightarrow{d} cX$

- ❸ $X_n/Y_n \xrightarrow{d} X/c$ as long as $c \neq 0$

- These results follow from CMT, plus the fact that

$$(X_n, Y_n) \xrightarrow{d} (X, c)$$

provided that c is non-random.

Delta method

- If μ is non-random,

$$\sqrt{n}(X_n - \mu) \xrightarrow{d} Z$$

and g is differentiable at μ , then

$$\sqrt{n}(g(X_n) - g(\mu)) \xrightarrow{d} g'(\mu)Z$$

- Remarks.

- 1 No assumption about normality is needed for Z .
- 2 The theorem goes through in the multivariate case

★ Example: $F(\gamma_1, \gamma_2) = \gamma_1\gamma_2$

Results about averages

- Let $\{X_i\}$ be sampled iid from some distribution. Denote

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

Results about averages

- Let $\{X_i\}$ be sampled iid from some distribution. Denote

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

- Law of large numbers. If $\{X_i\}$ are sampled iid from distribution with finite mean, then

$$\bar{X}_n \xrightarrow{p} \mu$$

Results about averages

- Let $\{X_i\}$ be sampled iid from some distribution. Denote

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

- Law of large numbers. If $\{X_i\}$ are sampled iid from distribution with finite mean, then

$$\bar{X}_n \xrightarrow{p} X_1$$

- Central limit theorem. If $\{X_i\}$ are sampled iid from distribution with finite mean μ and variance σ^2 ,

$$\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} Z \sim N(0, \sigma^2)$$

Results about averages

- **Remark.** If X_i have finite mean and variance, the central limit theorem gives a bound on the rate of convergence of \bar{X}_n

Results about averages

- **Remark.** If X_i have finite mean and variance, the central limit theorem gives a bound on the rate of convergence of \bar{X}_n
- In fact:

$$\bar{X}_n - \mu = \frac{1}{\sqrt{n}} [\sqrt{n}(\bar{X}_n - \mu)] = O_p(n^{-1/2})O_p(1) = O_p(n^{-1/2})$$

Results about averages

- **Remark.** If X_i have finite mean and variance, the central limit theorem gives a bound on the rate of convergence of \bar{X}_n
- In fact:

$$\bar{X}_n - \mu = \frac{1}{\sqrt{n}} [\sqrt{n}(\bar{X}_n - \mu)] = O_p(n^{-1/2})O_p(1) = O_p(n^{-1/2})$$

- Hence (with finite second moments) the convergence of \bar{X}_n to μ is no slower than $n^{-1/2}$ goes to 0.

OLS asymptotics: the linear projection model

- Consider the following assumptions:

(1) (y_i, x_i) sampled iid

(2) $(y_i^2) < \infty$

(2') $(y_i^4) < \infty$

(3) $(\|x_i\|^2) < \infty$

(3') $(\|x_i\|^4) < \infty$

(4) $[x_i x_i']$ positive definite

- Notation:

$$[x_i x_i'] =: Q_{xx}$$

$$[x_i y_i] =: Q_{xy}$$

OLS asymptotics: the linear projection model

- Results:

- Under (1), (2), (3) and (4), OLS is consistent for the population projection coefficient:

$$\hat{\beta} = \left(\sum_{i=1}^n x_i x_i' \right)^{-1} \sum_{i=1}^n x_i y_i \xrightarrow{P} \beta := Q_{xx}^{-1} Q_{xy}$$

- Under (1), (2'), (3') and (4), OLS is asymptotically normal:

$$\sqrt{n} \left[\left(\sum_{i=1}^n x_i x_i' \right)^{-1} \sum_{i=1}^n x_i y_i - \beta \right] \xrightarrow{d} N(0, Q_{xx}^{-1} \Omega Q_{xx}^{-1})$$

- ★ β is again the population projection coefficient

- ★ $\Omega = [u_i^2 x_i x_i']$

- ★ $u_i := y_i - x_i' \beta$

- Notation:

$$V_{\beta} = Q_{xx}^{-1} \Omega Q_{xx}^{-1}$$

OLS asymptotics: the linear regression model

- Consider the following assumptions:

(1) (y_i, x_i) sampled iid

(2) $(y_i^2) < \infty$

(3) $(\|x_i\|^2) < \infty$

(4) $[x_i x_i']$ positive definite

(5) $(y_i | x_i) = x_i' \beta$

- Assumptions (1), (2), (3), (4) and (5) are called the linear regression model

OLS asymptotics: the linear regression model

- Consider the following assumptions:

(1) (y_i, x_i) sampled iid

(2') $(y_i^4) < \infty$

(3') $(\|x_i\|^4) < \infty$

(4) $[x_i x_i']$ positive definite

(5) $(y_i | x_i) = x_i' \beta$

- Assumptions (1), (2), (3), (4) and (5) are called the linear regression model
- Again (2') and (3') are required for asymptotic normality

OLS asymptotics: the linear regression model

- Consider the following assumptions:

(1) (y_i, x_i) sampled iid

(2) $(y_i^2) < \infty$

(2') $(y_i^4) < \infty$

(3) $(\|x_i\|^2) < \infty$

(3') $(\|x_i\|^4) < \infty$

(4) $[x_i x_i']$ positive definite

(5) $(y_i|x_i) = x_i' \beta$

(6) $\mathbf{V}[y_i|x_i] \equiv \sigma^2$

- Assumptions (1), (2), (3), (4) and (5) are called the linear regression model
- Again (2') and (3') are required for asymptotic normality
- If (6) is assumed, we have the **homoskedastic** linear regression model

OLS asymptotics: the linear regression model

- Finite sample variance of OLS:

$$\hat{\beta} = (X'X)^{-1}X'Y \implies \mathbf{V}(\hat{\beta}|X) = (X'X)^{-1}X'\mathbf{V}(Y|X)X(X'X)^{-1}$$

OLS asymptotics: the linear regression model

- Finite sample variance of OLS:

$$\hat{\beta} = (X'X)^{-1}X'Y \implies \mathbf{V}(\hat{\beta}|X) = (X'X)^{-1}X'\mathbf{V}(Y|X)X(X'X)^{-1}$$

- $\mathbf{V}(Y|X)$ is a diagonal matrix. Why?

OLS asymptotics: the linear regression model

- Finite sample variance of OLS:

$$\hat{\beta} = (X'X)^{-1}X'Y \implies \mathbf{V}(\hat{\beta}|X) = (X'X)^{-1}X'\mathbf{V}(Y|X)X(X'X)^{-1}$$

- $\mathbf{V}(Y|X)$ is a diagonal matrix. Why?
- Denote $D = \mathbf{V}(Y|X)$. Hansen defines:

$$V_{\hat{\beta}} = (X'X)^{-1}X'DX(X'X)^{-1}$$

OLS asymptotics: the linear regression model

- Finite sample variance of OLS:

$$\hat{\beta} = (X'X)^{-1}X'Y \implies \mathbf{V}(\hat{\beta}|X) = (X'X)^{-1}X'\mathbf{V}(Y|X)X(X'X)^{-1}$$

- $\mathbf{V}(Y|X)$ is a diagonal matrix. Why?
- Denote $D = \mathbf{V}(Y|X)$. Hansen defines:

$$V_{\hat{\beta}} = (X'X)^{-1}X'DX(X'X)^{-1}$$

- This is **not** the same thing as

$$V_{\beta} = Q_{xx}^{-1}[u_i^2 x_i x_i']Q_{xx}^{-1}$$

OLS asymptotics: the linear regression model

- Finite sample variance of OLS:

$$\hat{\beta} = (X'X)^{-1}X'Y \implies \mathbf{V}(\hat{\beta}|X) = (X'X)^{-1}X'\mathbf{V}(Y|X)X(X'X)^{-1}$$

- $\mathbf{V}(Y|X)$ is a diagonal matrix. Why?
- Denote $D = \mathbf{V}(Y|X)$. Hansen defines:

$$V_{\hat{\beta}} = (X'X)^{-1}X'DX(X'X)^{-1}$$

- This is **not** the same thing as

$$V_{\beta} = Q_{xx}^{-1}[u_i^2 x_i x_i']Q_{xx}^{-1}$$

- It is true however that

$$nV_{\hat{\beta}} \xrightarrow{p} V_{\beta}$$

(why?)

OLS asymptotics: estimating the limiting variance

- We know that the limiting variance of OLS in the LPM is

$$V_{\beta} = Q_{xx}^{-1} \Omega Q_{xx}^{-1}$$

- A valid estimator for $\Omega = [u_i^2 x_i x_i']$ is

$$\hat{\Omega} = \frac{1}{n} \sum_{i=1}^n \hat{u}_i^2 x_i x_i'$$

- One estimator for V_{β} is

$$\hat{V}_{\beta}^{HC0} = \hat{Q}_{xx}^{-1} \hat{\Omega} \hat{Q}_{xx}^{-1}$$

Prediction: regression intervals

- One object of interest in the LRM is

$$m(x) = [y_i|x] = x'\beta$$

Prediction: regression intervals

- One object of interest in the LRM is

$$m(x) = [y_i|x] = x'\beta$$

- How does one estimate this function? Reasonable candidate is:

$$\hat{m}(x) = x'\hat{\beta}$$

Prediction: regression intervals

- One object of interest in the LRM is

$$m(x) = [y_i|x] = x'\beta$$

- How does one estimate this function? Reasonable candidate is:

$$\hat{m}(x) = x'\hat{\beta}$$

- Note that for fixed x ,

$$\sqrt{n}(\hat{m}(x) - m(x)) = x'\sqrt{n}(\hat{\beta} - \beta) \xrightarrow{d} N(0, x'V_{\beta}x)$$

so that

$$\left(x'\hat{V}_{\beta}x\right)^{-1/2} \sqrt{n}[\hat{m}(x) - m(x)] \xrightarrow{d} N(0, 1)$$

Prediction: regression intervals

- One object of interest in the LRM is

$$m(x) = [y_i|x] = x'\beta$$

- How does one estimate this function? Reasonable candidate is:

$$\hat{m}(x) = x'\hat{\beta}$$

- Note that for fixed x ,

$$\sqrt{n}(\hat{m}(x) - m(x)) = x'\sqrt{n}(\hat{\beta} - \beta) \xrightarrow{d} N(0, x'V_{\beta}x)$$

so that

$$\left(x'\hat{V}_{\beta}x\right)^{-1/2} \sqrt{n}[\hat{m}(x) - m(x)] \xrightarrow{d} N(0, 1)$$

- That motivates the approximation

$$\hat{m}(x) \approx N\left(m(x), \frac{1}{n}x'\hat{V}_{\beta}x\right)$$

Prediction: forecast errors

- Suppose you expect to obtain one additional data point: (y_{n+1}, x_{n+1})

Prediction: forecast errors

- Suppose you expect to obtain one additional data point: (y_{n+1}, x_{n+1})
- How would you predict y_{n+1} ? Quadratic loss function implies best prediction is

$$y_{n+1}^* = [y_i | x_{n+1}] = m(x_{n+1}) = x_{n+1}'\beta$$

Prediction: forecast errors

- Suppose you expect to obtain one additional data point: (y_{n+1}, x_{n+1})
- How would you predict y_{n+1} ? Quadratic loss function implies best prediction is

$$y_{n+1}^* = [y_i | x_{n+1}] = m(x_{n+1}) = x'_{n+1} \beta$$

- The above is infeasible; use plug-in estimate $\hat{m}(x)$

$$\tilde{y}_{n+1} = x'_{n+1} \hat{\beta}$$

Prediction: forecast errors

- Suppose you expect to obtain one additional data point: (y_{n+1}, x_{n+1})
- How would you predict y_{n+1} ? Quadratic loss function implies best prediction is

$$y_{n+1}^* = [y_i | x_{n+1}] = m(x_{n+1}) = x'_{n+1} \beta$$

- The above is infeasible; use plug-in estimate $\hat{m}(x)$

$$\tilde{y}_{n+1} = x'_{n+1} \hat{\beta}$$

- Forecast error is

$$\tilde{e}_{n+1} = y_{n+1} - x' \hat{\beta} = x' \beta + e_i - x' \hat{\beta} = e_{n+1} + x'(\beta - \hat{\beta})$$

- ▶ The true regression has an error e_{n+1}
- ▶ In addition to that, there is an estimation noise coming from $\hat{\beta}$ being different than β

Prediction: forecast errors

- Finite sample variance of e_{n+1} :

$$[\hat{e}_{n+1}^2 | x_{n+1}] = \sigma^2(x_{n+1}) + x'_{n+1} V_{\hat{\beta}} x_{n+1}$$

- Can we build a confidence interval for the forecast?